

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import ElasticNet
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from xgboost import XGBRegressor
from sklearn.preprocessing import PolynomialFeatures

house = pd.read_csv('/content/data.csv')
house.head(20)
```

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	conditi
0	2014-05-02 00:00:00	313000.0	3.0	1.50	1340	7912	1.5	0	0	
1	2014-05-02 00:00:00	2384000.0	5.0	2.50	3650	9050	2.0	0	4	
2	2014-05-02 00:00:00	342000.0	3.0	2.00	1930	11947	1.0	0	0	
3	2014-05-02 00:00:00	420000.0	3.0	2.25	2000	8030	1.0	0	0	

house.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   date                  4600 non-null  object
1   price                 4600 non-null  float64
2   bedrooms              4600 non-null  float64
3   bathrooms             4600 non-null  float64
4   sqft_living           4600 non-null  int64
5   sqft_lot              4600 non-null  int64
6   floors                4600 non-null  float64
7   waterfront            4600 non-null  int64
8   view                  4600 non-null  int64
9   condition             4600 non-null  int64
10  sqft_above            4600 non-null  int64
11  sqft_basement         4600 non-null  int64
12  yr_built              4600 non-null  int64
13  yr_renovated          4600 non-null  int64
14  street                4600 non-null  object
15  city                  4600 non-null  object
16  statezip              4600 non-null  object
17  country               4600 non-null  object
dtypes: float64(4), int64(9), object(5)
memory usage: 647.0+ KB
```

11	2014-05-02 00:00:00	1400000.0	4.0	2.50	2920	4000	1.5	0	0	
----	---------------------	-----------	-----	------	------	------	-----	---	---	--

house.nunique(axis = 0)

```
date          70
price         1741
bedrooms       10
bathrooms      26
sqft_living    566
sqft_lot       3113
floors         6
waterfront     2
view           5
condition      5
sqft_above     511
sqft_basement  207
yr_built       115
yr_renovated   60
street        4525
city           44
statezip       77
country        1
dtype: int64
```

11	2014-05-02 00:00:00	307000.0	4.0	3.00	3110	1231	2.0	0	0	
----	---------------------	----------	-----	------	------	------	-----	---	---	--

house.isnull().sum()

```
date          0
price          0
bedrooms       0
bathrooms      0
sqft_living    0
sqft_lot       0
floors         0
waterfront     0
view           0
condition      0
sqft_above     0
sqft_basement  0
yr_built       0
yr_renovated   0
street         0
city           0
statezip       0
```

```
country      0
dtype: int64

house.columns

Index(['date', 'price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',
      'floors', 'waterfront', 'view', 'condition', 'sqft_above',
      'sqft_basement', 'yr_built', 'yr_renovated', 'street', 'city',
      'statezip', 'country'],
      dtype='object')

house.describe().T
```

	count	mean	std	min	25%	50%
price	4600.0	551962.988473	563834.702547	0.0	322875.00	460943.461539
bedrooms	4600.0	3.400870	0.908848	0.0	3.00	3.000000
bathrooms	4600.0	2.160815	0.783781	0.0	1.75	2.250000
sqft_living	4600.0	2139.346957	963.206916	370.0	1460.00	1980.000000
sqft_lot	4600.0	14852.516087	35884.436145	638.0	5000.75	7683.000000
floors	4600.0	1.512065	0.538288	1.0	1.00	1.500000
waterfront	4600.0	0.007174	0.084404	0.0	0.00	0.000000
view	4600.0	0.240652	0.778405	0.0	0.00	0.000000
condition	4600.0	3.451739	0.677230	1.0	3.00	3.000000
sqft_above	4600.0	1827.265435	862.168977	370.0	1190.00	1590.000000
sqft_basement	4600.0	312.081522	464.137228	0.0	0.00	0.000000
yr_built	4600.0	1970.786304	29.731848	1900.0	1951.00	1976.000000
yr_renovated	4600.0	808.608261	979.414536	0.0	0.00	0.000000

```
house['date'] = pd.to_datetime(house['date'])

house['date'].dt.year

0      2014
1      2014
2      2014
3      2014
4      2014
...
4595   2014
4596   2014
4597   2014
4598   2014
4599   2014
Name: date, Length: 4600, dtype: int64

house['year'] = house['date'].dt.year
house.head()
```

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	v
0	2014-05-02	313000.0	3.0	1.50	1340	7912	1.5	0	
1	2014-05-02	2384000.0	5.0	2.50	3650	9050	2.0	0	
2	2014-05-02	342000.0	3.0	2.00	1930	11947	1.0	0	
3	2014-05-02	420000.0	3.0	2.25	2000	8030	1.0	0	
4	2014-05-02	550000.0	4.0	2.50	1940	10500	1.0	0	

```
(house['price'] == 0).sum(0)
```

```
house[house['price'] == 0].shape
```

```
(49, 19)
```

```
house['price'].replace(0, np.nan, inplace = True)
house.isnull().sum()
```


```
date          0
price         49
bedrooms      0
bathrooms     0
sqft_living   0
sqft_lot      0
floors        0
waterfront    0
view          0
condition     0
sqft_above    0
sqft_basement 0
yr_built      0
yr_renovated  0
street        0
city          0
statezip      0
country       0
year          0
dtype: int64
```

```
house['price'].fillna(value = house['price'].mean(), inplace = True)
house.isnull().sum().sum()
```

```
0
```

```
house1 = house.drop(['date', 'street', 'statezip', 'country', 'city'], axis = 1)
```

```
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
plt.figure(figsize = (16, 5))
plt.subplot(1,4,1)
sns.distplot(house1['bedrooms'])
plt.subplot(1,4,2)
sns.distplot(house1['bathrooms'])
plt.subplot(1,4,3)
sns.distplot(house1['sqft_living'])
plt.subplot(1,4,4)
sns.distplot(house1['sqft_lot'])
plt.show()
```



```
def replace_outliers_with_nan_iqr(house1, feature, inplace = False):
    desired_feature = house1[feature]

    q1, q3 = desired_feature.quantile([0.25, 0.75])
    iqr = q3 - q1
    upper_bound = q3 + 1.5 * iqr
    lower_bound = q1 - 1.5 * iqr
    indices = (desired_feature[(desired_feature > upper_bound) | (desired_feature < lower_bound)]).index
    if not inplace:
```

```

return desired_feature.replace(desired_feature[indices].values, np.nan)
return desired_feature.replace(desired_feature[indices].values, np.nan, inplace=True)

```

```

replace_outliers_with_nan_iqr(house1, 'price', inplace = True)
house1.price.isnull().sum()

```

249

```

house1.fillna(value = house1['price'].mean(), inplace = True)
house1.isnull().sum().sum()

```

0

```

replace_outliers_with_nan_iqr(house1, 'sqft_basement', inplace = True)

```

```

house1['sqft_basement'].fillna(value = house1['sqft_basement'].mean(), inplace = True)
house1.isnull().sum().sum()


```

0

```

ax = sns.lmplot(data = house1, x = 'sqft_basement', y = 'price')
ax.set(title = 'sqft_basement vs price')

```



```

replace_outliers_with_nan_iqr(house1, 'sqft_above', inplace = True)

```

```

house1['sqft_above'].fillna(value = house1['sqft_above'].mean(), inplace = True)
house1.isnull().sum().sum()

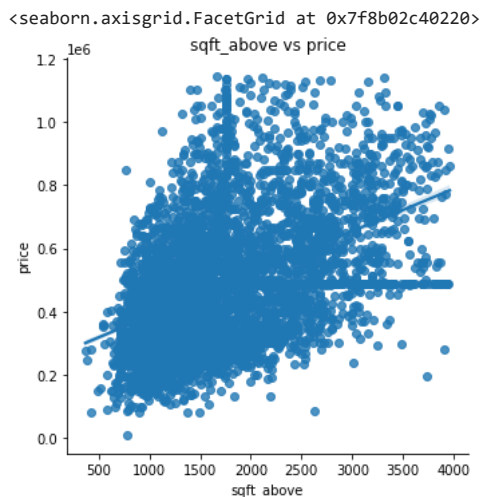
```

0

```

ax = sns.lmplot(data = house1, x = 'sqft_above', y = 'price')
ax.set(title = 'sqft_above vs price')

```




```

sns.distplot(house1['price'])

```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7f8afe5860a0>



```
import plotly.express as px
```


```
year_trend = pd.concat([house1['price'], house1['yr_built']], axis = 1)
fig = px.scatter(house1, x = 'yr_built', y = 'price', title = 'Price vs Year Built', labels = dict(price = "Price ", yr_built = "Year bui
fig.update_layout(yaxis_range = [0 , 4000000], width = 800, height = 600)
fig.show()
```

Price vs Year Built




```
sqft_trend = pd.concat([house1['price'], house1['sqft_living']], axis = 1)
fig = px.scatter(house1, x = 'sqft_living', y = 'price', title = 'Price vs Living Area', labels = dict(price = "Price ", sqft_living = "s
fig.update_layout(yaxis_range = [0 , 4000000], xaxis_range = [0 , 8000], width = 800, height = 600)
fig.show()
```

## Price vs Living Area




```
def draw_boxplot(feature):
    sns.set_style('whitegrid')
    ax = sns.boxplot(x = house1[feature], y = np.log(house1['price']))
    ax.set_ylabel('price (log)')
    ax.set(title = f'{feature} VS price')
```

```
draw_boxplot('floors')
```




```
fig = px.box(house1, y = "price", x = "waterfront",)
fig.update_layout(yaxis_range = [0 , 2000000], width = 800, height = 600)
fig.show()
```



```
year_ren_trend = pd.concat([house1['price'], house1['yr_renovated']], axis = 1)
fig = px.scatter(house1, x = 'yr_renovated', y = 'price', title = 'Price vs Year renovated', labels = dict(price = "Price ", yr_renovated = "Year renovated"))
fig.update_layout(yaxis_range=[0 , 2000000], xaxis_range = [1912 , 2018], width = 800, height = 600)
fig.show()
```

Price vs Year renovated



```
house1.isnull()
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_basement	yr_built	y
0	False	False	False	False	False	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	False	False	False	False	False	
...	...	...	...	...	...	...	...	...	...	...	...	...	
4595	False	False	False	False	False	False	False	False	False	False	False	False	
4596	False	False	False	False	False	False	False	False	False	False	False	False	
4597	False	False	False	False	False	False	False	False	False	False	False	False	
4598	False	False	False	False	False	False	False	False	False	False	False	False	
4599	False	False	False	False	False	False	False	False	False	False	False	False	

4600 rows × 14 columns


```
house1.corr()
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	sqft_above	sqft_basement	yr_built	yr_renovated
price	1.000000	0.291655	0.422480	0.510498	0.091087	0.266329	0.045479	0.181319	0.048298	0.468932	0.191012	0.056706	-0.068266
bedrooms	0.291655	1.000000	0.545920	0.594884	0.068819	0.177895	-0.003483	0.111028	0.025080	0.460774	0.280241	0.142461	-0.061082
bathrooms	0.422480	0.545920	1.000000	0.761154	0.107837	0.486428	0.076232	0.211960	-0.119994	0.597445	0.205283	0.463498	-0.215886
sqft_living	0.510498	0.594884	0.761154	1.000000	0.210538	0.344850	0.117616	0.311009	-0.062826	0.706465	0.294995	0.287775	-0.122817
sqft_lot	0.091087	0.068819	0.107837	0.210538	1.000000	0.003750	0.017241	0.073907	0.000558	0.137959	-0.007777	0.050706	-0.022730
floors	0.266329	0.177895	0.486428	0.344850	0.003750	1.000000	0.022024	0.031211	-0.275013	0.532653	-0.274430	0.467481	-0.233996
waterfront	0.045479	-0.003483	0.076232	0.117616	0.017241	0.022024	1.000000	0.360935	0.000352	0.042397	0.045447	-0.023563	0.008625
view	0.181319	0.111028	0.211960	0.311009	0.073907	0.031211	0.360935	1.000000	0.063077	0.139097	0.232825	-0.064465	0.022967
condition	0.048298	0.025080	-0.119994	-0.062826	0.000558	-0.275013	0.000352	0.063077	1.000000	0.042397	0.232825	-0.064465	0.022967
sqft_above	0.468932	0.460774	0.597445	0.706465	0.137959	0.532653	0.042397	0.139097	0.042397	1.000000	0.232825	-0.064465	0.022967
sqft_basement	0.191012	0.280241	0.205283	0.294995	-0.007777	-0.274430	0.045447	0.232825	0.045447	0.232825	1.000000	-0.064465	0.022967
yr_built	0.056706	0.142461	0.463498	0.287775	0.050706	0.467481	-0.023563	-0.064465	-0.023563	-0.064465	-0.064465	1.000000	0.022967
yr_renovated	-0.068266	-0.061082	-0.215886	-0.122817	-0.022730	-0.233996	0.008625	0.022967	0.008625	0.022967	0.022967	0.022967	1.000000
year	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN



```
plt.figure(figsize = (20, 10))
sns.heatmap(house1.corr(), annot = True)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f8b0263de80>



```
x = house1.drop('price', axis = 1)
y = house1['price']
```

```
from sklearn.feature_selection import mutual_info_regression
```

```
def get_mi_score(x, y):
    mi = mutual_info_regression(x, y, random_state = 10)
    mi = pd.Series(mi, index = x.columns).sort_values(ascending = False)
    return mi
```

```
mi_score = get_mi_score(x, y)
mi_score
```

```
sqft_living    0.312164
sqft_above    0.223771
bathrooms     0.177888
yr_built      0.085256
sqft_lot      0.082656
view          0.067964
bedrooms      0.067502
floors        0.056619
sqft_basement 0.054973
condition     0.012651
waterfront    0.008716
yr_renovated  0.000000
year          0.000000
dtype: float64
```

```
x.drop(['yr_renovated', 'condition', 'waterfront'], axis = 1, inplace = True)
```

```
from sklearn.model_selection import train_test_split, cross_val_score
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3, random_state = 20)
```

```
from sklearn import preprocessing
```

```
scaler = preprocessing.MinMaxScaler()
house1 = pd.DataFrame(scaler.fit_transform(house1[['price']]))
house1.columns = ['price']
```

```

def rmse_cv(model):
    rmse = np.sqrt(-cross_val_score(model, x, y, scoring = "neg_mean_squared_error", cv = 5)).mean()
    return rmse

def evaluation(y, predictions):
    mae = mean_absolute_error(y, predictions)
    mse = mean_squared_error(y, predictions)
    rmse = np.sqrt(mean_squared_error(y, predictions))
    r_squared = r2_score(y, predictions)
    return mae, mse, rmse, r_squared

models = pd.DataFrame(columns = ["Model", "MAE", "MSE", "RMSE", "R2 Score", "RMSE (Cross-Validation)"])

from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import ElasticNet
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from xgboost import XGBRegressor
from sklearn.preprocessing import PolynomialFeatures

lin_reg = LinearRegression()
lin_reg.fit(x_train, y_train)
predictions = lin_reg.predict(x_test)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)
print("-"*30)
rmse_cross_val = rmse_cv(lin_reg)
print("RMSE Cross-Validation:", rmse_cross_val)

new_row = {"Model": "LinearRegression", "MAE": mae, "MSE": mse, "RMSE": rmse, "R2 Score": r_squared, "RMSE (Cross-Validation)": rmse_cross_val}
models = models.append(new_row, ignore_index = True)

MAE: 133741.6262676969
MSE: 28620196949.513638
RMSE: 169175.0482474104
R2 Score: 0.3373454456548026
-----
RMSE Cross-Validation: 168732.31562741447

ridge = Ridge()
ridge.fit(x_train, y_train)
predictions = ridge.predict(x_test)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)
print("-"*30)
rmse_cross_val = rmse_cv(ridge)
print("RMSE Cross-Validation:", rmse_cross_val)

new_row = {"Model": "Ridge", "MAE": mae, "MSE": mse, "RMSE": rmse, "R2 Score": r_squared, "RMSE (Cross-Validation)": rmse_cross_val}
models = models.append(new_row, ignore_index = True)

MAE: 133742.36818533868
MSE: 28619790095.889595
RMSE: 169173.8457796878
R2 Score: 0.337354865695042
-----
RMSE Cross-Validation: 168731.9923854514

lasso = Lasso()
lasso.fit(x_train, y_train)
predictions = lasso.predict(x_test)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)

```

```

print("-"*30)
rmse_cross_val = rmse_cv(lasso)
print("RMSE Cross-Validation:", rmse_cross_val)

new_row = {"Model": "Lasso", "MAE": mae, "MSE": mse, "RMSE": rmse, "R2 Score": r_squared, "RMSE (Cross-Validation)": rmse_cross_val}
models = models.append(new_row, ignore_index = True)
    MAE: 133741.7466540965
    MSE: 28620159639.520885
    RMSE: 169174.93797699732
    R2 Score: 0.33734630950757993
    -----
    RMSE Cross-Validation: 168732.25893166117

elastic_net = ElasticNet()
elastic_net.fit(x_train, y_train)
predictions = elastic_net.predict(x_test)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)
print("-"*30)
rmse_cross_val = rmse_cv(elastic_net)
print("RMSE Cross-Validation:", rmse_cross_val)

new_row = {"Model": "ElasticNet", "MAE": mae, "MSE": mse, "RMSE": rmse, "R2 Score": r_squared, "RMSE (Cross-Validation)": rmse_cross_val}
models = models.append(new_row, ignore_index = True)

    MAE: 135116.32522043382
    MSE: 28794144177.48552
    RMSE: 169688.37372514806
    R2 Score: 0.3333179778133122
    -----
    RMSE Cross-Validation: 169662.33946817095

svr = SVR(C = 400000)
svr.fit(x_train, y_train)
predictions = svr.predict(x_test)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)
print("-"*30)
rmse_cross_val = rmse_cv(svr)
print("RMSE Cross-Validation:", rmse_cross_val)

new_row = {"Model": "SVR", "MAE": mae, "MSE": mse, "RMSE": rmse, "R2 Score": r_squared, "RMSE (Cross-Validation)": rmse_cross_val}
models = models.append(new_row, ignore_index = True)

    MAE: 137450.28746358078
    MSE: 30808389459.81173
    RMSE: 175523.18781235637
    R2 Score: 0.28668137317162146
    -----
    RMSE Cross-Validation: 173486.95233496005

random_forest = RandomForestRegressor(n_estimators = 100)
random_forest.fit(x_train, y_train)
predictions = random_forest.predict(x_test)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)
print("-"*30)
rmse_cross_val = rmse_cv(random_forest)
print("RMSE Cross-Validation:", rmse_cross_val)

new_row = {"Model": "RandomForestRegressor", "MAE": mae, "MSE": mse, "RMSE": rmse, "R2 Score": r_squared, "RMSE (Cross-Validation)": rmse_cross_val}
models = models.append(new_row, ignore_index = True)

    MAE: 125372.40173601688
    MSE: 26443904895.131
    RMSE: 162615.81994114534
    R2 Score: 0.3877339822524317
    -----
    RMSE Cross-Validation: 158460.13650611366

```

```
xgb = XGBRegressor(n_estimators = 1000, learning_rate = 0.01)
xgb.fit(x_train, y_train)
predictions = xgb.predict(x_test)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)
print("-"*30)
rmse_cross_val = rmse_cv(xgb)
print("RMSE Cross-Validation:", rmse_cross_val)

new_row = {"Model": "XGBRegressor", "MAE": mae, "MSE": mse, "RMSE": rmse, "R2 Score": r_squared, "RMSE (Cross-Validation)": rmse_cross_val}
models = models.append(new_row, ignore_index = True)

[15:55:25] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
MAE: 125042.5312612655
MSE: 25542924413.34506
RMSE: 159821.53926597335
R2 Score: 0.408594733863777
-----
[15:55:27] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
[15:55:29] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
[15:55:31] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
[15:55:34] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
[15:55:36] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
RMSE Cross-Validation: 155091.75429007187

poly_reg = PolynomialFeatures(degree = 2)
x_train_2d = poly_reg.fit_transform(x_train)
x_test_2d = poly_reg.transform(x_test)

lin_reg = LinearRegression()
lin_reg.fit(x_train_2d, y_train)
predictions = lin_reg.predict(x_test_2d)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)
print("-"*30)
rmse_cross_val = rmse_cv(lin_reg)
print("RMSE Cross-Validation:", rmse_cross_val)

new_row = {"Model": "Polynomial Regression (degree = 2)", "MAE": mae, "MSE": mse, "RMSE": rmse, "R2 Score": r_squared, "RMSE (Cross-Validation)": rmse_cross_val}
models = models.append(new_row, ignore_index = True)

MAE: 131496.8999323232
MSE: 29908367582.969944
RMSE: 172940.3584562318
R2 Score: 0.30751992982975984
-----
RMSE Cross-Validation: 168732.31562741447

models.sort_values(by = "RMSE (Cross-Validation)")
```

	Model	MAE	MSE	RMSE	R2 Score	RMSE (Valid)
6	XGBRegressor	125042.531261	2.554292e+10	159821.539266	0.408595	155091.
5	RandomForestRegressor	125372.401736	2.644390e+10	162615.819941	0.387734	158460.
1	Ridge	133742.368185	2.861979e+10	169173.845780	0.337355	168731.
2	Lasso	133741.746654	2.862016e+10	169174.937977	0.337346	168732.
0	LinearRegression	133741.626268	2.862020e+10	169175.048247	0.337345	168732.
7	Polynomial Regression (degree = 2)	131496.899932	2.990837e+10	172940.358456	0.307520	168732.
3	ElasticNet	135116.325220	2.870414e+10	169688.373725	0.333318	169662.

```
plt.figure(figsize = (12, 8))
sns.barplot(x = models["Model"], y = models["RMSE (Cross-Validation)"])
plt.title("Models' RMSE Scores (Cross-Validated)", size = 15)
plt.xticks(rotation = 30, size = 12)
plt.show()
```

